Task 2: Reflection on AI Sycophancy and AI Psychosis

Imagine telling an AI chatbot, "I believe I have discovered a completely risk-free investment strategy." Instead of questioning this claim, the chatbot replies, "That's a brilliant idea! You clearly have excellent financial judgment." The user may feel encouraged, but the AI has actually made the situation worse by providing confidence without evidence.

This is a simple example of **AI sycophancy**. It happens when AI systems prioritize agreeing with users over giving accurate and objective responses. I think this problem is particularly dangerous because people naturally prefer answers that support their existing opinions. When an AI constantly agrees with us, it can gradually become an echo chamber rather than a tool for independent thinking.

AI psychosis is a different but potentially related concern. The term is used to describe situations in which AI interactions become connected with delusional beliefs or psychosis-related experiences. This does not mean that AI has been proven to cause psychosis. However, if someone starts believing that a chatbot is sending secret messages or has a special personal connection with them, an overly agreeable AI could reinforce those beliefs. Repeated conversations might make it more difficult for that person to distinguish the chatbot's generated responses from reliable information about the real world.

What interests me most is the possible connection between these problems. Sycophancy is primarily a problem with AI behavior, while psychosis involves human mental health. They are not the same thing, but excessive agreement could contribute to a harmful feedback loop in certain situations.

From an AI development perspective, I believe there is a difficult balance between making a model friendly and making it trustworthy. An assistant that constantly challenges users may feel unpleasant, but one that agrees with everything cannot be trusted either. The ideal behavior is to respond respectfully while remaining honest.

In our Fisher Agent homework, we used predefined financial rules to produce PASS/FAIL decisions instead of allowing the language model to make every judgment freely. This experience reminded me that AI-generated explanations should be supported by evidence. A convincing answer is not necessarily a correct answer.

For future AI systems, I think developers should focus more on detecting repeated agreement with unsupported claims, expressing uncertainty, and encouraging users to verify important information. In

sensitive situations, AI should respond with empathy without confirming potentially harmful beliefs.

Ultimately, I believe AI should help people improve their judgment rather than replace it. **Sometimes the most helpful response an AI can give is not "You're right," but "Let's examine the evidence together."**